

## HS650 Homework Header

HW 3

Winter 2018, DSPA (HS650)

Name: Xinchun Li

SID: 85922126

UMich E-mail: xincli@umich.edu

I certify that the following paper represents my own independent work and conforms with the guidelines of academic honesty described in the UMich student handbook.

### Problem 3.1 (Probability Distributions):

Complete the following tasks for each of the probability distributions below:

- Generate plots of the [density, CDF, and the quantile \(inverse-CDF\) functions \(Links to an external site.\)](#)Links to an external site.
- Report the first [4 moments \(mean, variance, skewness, kurtosis\) \(Links to an external site.\)](#)Links to an external site.
- Complete the discrete probability distributions table below. The cell values in the table represent the values of the quantile function for the corresponding p-value (column) and distribution (row).

**# Compute the values of the quantile function for the corresponding p-value and distribution**

```
>s <- seq(0.1, 0.9, 0.1)
>sweibull <- qweibull(s, 1, 5)
>sunif <- qunif(s, 1, 10)
>sqt <- qt(s, 1)
>scauchy <- qcauchy(s)
>sbinom <- qnbinom(s, 10, 0.5)
>schisq <- qchisq(s, 10)
```

```
>stable <- cbind(sweibull, sunif, sqrt, scauchy, sbinom, schisq)
>t(stable)
```

### # Output

	sweibull	sunif	sqrt	scauchy	sbinom	schisq
[1,]	0.5268026	1.9	-3.0776835	-3.0776835	5	4.865182
[2,]	1.1157178	2.8	-1.3763819	-1.3763819	6	6.179079
[3,]	1.7833747	3.7	-0.7265425	-0.7265425	7	7.267218
[4,]	2.5541281	4.6	-0.3249197	-0.3249197	8	8.295472
[5,]	3.4657359	5.5	0.0000000	0.0000000	9	9.341818
[6,]	4.5814537	6.4	0.3249197	0.3249197	11	10.473236
[7,]	6.0198640	7.3	0.7265425	0.7265425	12	11.780723
[8,]	8.0471896	8.2	1.3763819	1.3763819	14	13.441958
[9,]	11.5129255	9.1	3.0776835	3.0776835	16	15.987179

### # Distribution Table

```
>smatrix <- as.matrix(t(stable))

>stable1 <- knitr::kable(smatrix, digits = 4, col.names = c("0.1",
"0.2", "0.3", "0.4", "0.5", "0.6", "0.7", "0.8", "0.9"))

> dnorm(s)

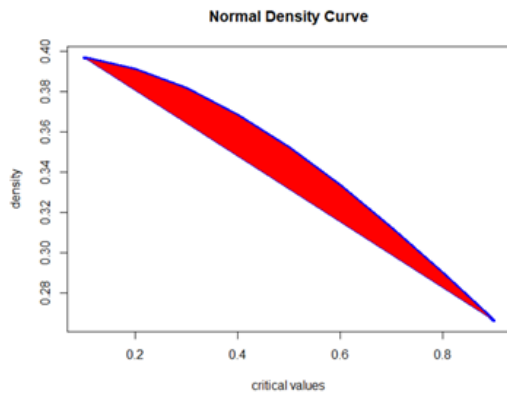
[1] 0.3969525 0.3910427 0.3813878 0.3682701 0.3520653
0.3332246 0.3122539 0.2896916 0.2660852
```

### # Density Plot

```
>dStandardNormal <- data.frame(Z=s, Density=dnorm(s, mean=0,
sd=1), Distribution=pnorm(s, mean=0, sd=1))

>plot(s, dStandardNormal$Density, main="Normal Density Curve",
type = "l", xlab = "critical values", ylab="density", lwd=4,
col="blue")


>polygon(s, dStandardNormal$Density, col="red", border="blue")
```



## # Cumulative Distribution Function Plot

```
>plot(s, dStandardNormal$Distribution, main="Normal
Distribution", type = "l", xlab = "critical values", ylab="Cumulative
Distribution", lwd=4, col="blue")
```

```
>polygon(s, dStandardNormal$Distribution, col="red",
border="blue")
```

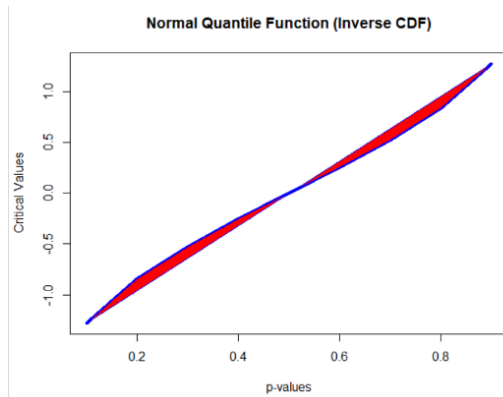


## # The Quantile Function Plot

```
>qStandardNormal <- data.frame(Q=s, Quantile=qnorm(s,
mean=0, sd=1))
```

```
>plot(s, qStandardNormal$Quantile, main="Normal Quantile
Function (Inverse CDF)", type = "l", xlab = "p-values",
ylab="Critical Values", lwd=4, col="blue")
```

```
>polygon(s, qStandardNormal$Quantile, col="red", border="blue")
```



### # Report Four Moments (mean, variance, skewness, kurtosis) for each set of data in a table

```
>weibull_moments <- c(mean(sweibull), var(sweibull),
skewness(sweibull), kurtosis(sweibull))

>unif_moments <- c(mean(sunif), var(sunif), skewness(sunif),
kurtosis(sunif))

>qt_moments <- c(mean(sqt), var(sqt), skewness(sqt),
kurtosis(sqt))

>cauchy_moments <- c(mean(scauchy), var(scauchy),
skewness(scauchy), kurtosis(scauchy))

>binom_moments <- c(mean(sbinom), var(sbinom),
skewness(sbinom), kurtosis(sbinom))

>chisq_moments <- c(mean(schisq), var(schisq),
skewness(schisq), kurtosis(schisq))

>four_moments <- c(weibull_moments, unif_moments,
qt_moments, cauchy_moments, binom_moments, chisq_moments)

>four_moments

[1] 4.400799e+00 1.293345e+01 7.019402e-01 -9.121017e-
01 5.500000e+00 6.075000e+00 -2.397404e-16 -1.601481e+00

[9] 1.356457e-16 3.000000e+00 2.845866e-16 -6.888889e-01
1.356457e-16 3.000000e+00 2.845866e-16 -6.888889e-01

[17] 9.777778e+00 1.394444e+01 2.954668e-01 -
1.492591e+00 9.736874e+00 1.277971e+01 3.044958e-01 -
1.300956e+00
```

### Problem 3.2 (Matrix equation solution):

**Solve the following system of linear equations and validate your solution. Validate your solution.**

$$6x + 3y - 3z + w = 2$$

$$7x + y + 2z + 2w = 5$$

$$5x + 3y - 3z + w = 3$$

$$-6x - 2y + 3z = 6$$

```
A_matrix_values <- c(6, 3, -3, 1, 7, 1, 2, 2, 5, 3, -3, 1, -6, -2, 3, 0)
```

```
A <- matrix(A_matrix_values, nrow = 4, ncol = 4)
```

```
b <- c(2, 5, 3, 6)
```

```
# to solve Ax = b, x = A ^ {-1} * b
```

```
x <- solve(A, b)
```

```
# Ax = b ==> x = A ^ {-1} * b
```

```
x
```

```
[1] -25  3  25 -1
```

### Verification:

```
> A.inverse <- solve(A) # the inverse matrix A^{-1}
```

```
> x1 <- A.inverse %*% b
```

```
> x1
```

```
      [,1]
```

```
[1,] -25
```

```
[2,]  3
```

```
[3,] 25
```

```
[4,] -1
```

**As we can see, the x and x1 are the same.**

### Problem 3.3 (Dimensionality reduction)

**Use PCA and t-SNE to analyze and interpret the monthly US Federal Reserve Monetary-Base Data (1959-2009)**[Links to an external site.](#)

#### # PCA

```
> ReserveData.sub <- ReserveData[, -1]
```

**We need to center the ReserveData.sub by subtracting the average of all column means from each element in the column. Next, we cast ReserveData.sub as a matrix and compute its variance covariance matrix, S. Finally, we can calculate the corresponding eigenvalues and eigenvectors of S.**

```
> mu <- apply(ReserveData.sub, 2, mean)
```

```
> mean(mu)
```

```
[1] 729.8481
```

```
> ReserveData.center <- as.matrix(ReserveData.sub)-mean(mu)
```

```
> S <- cov(ReserveData.center)
```

```
> eigen(S)
```

eigen() decomposition

\$values

```
[1] 6.459371e+06 8.105094e+04 1.954341e+04 5.516879e+03  
1.436806e+03 7.460492e+02 1.106305e+00
```

\$vectors

```
          [,1]      [,2]      [,3]      [,4]      [,5]  
[,6]      [,7]
```

```
[1,] 0.43802915 0.840194054 0.15824598 -0.22991060 -
0.154725598 0.01831078 0.005333293

[2,] 0.86861886 -0.366656200 0.02497568 0.32952256
0.008871566 0.04212471 -0.001550690

[3,] 0.17100350 -0.373093272 -0.03754267 -0.84575739 -
0.331941670 -0.06604044 -0.017250689

[4,] 0.11643248 0.097111962 -0.60687967 -0.17742808
0.488043646 0.06809777 -0.578277773

[5,] 0.01957359 0.071528131 -0.64779682 0.02767588 -
0.186822349 0.44733754 0.582306156

[6,] 0.01238692 0.076528374 -0.42823496 0.19664207 -
0.373796867 -0.79510424 -0.004673683

[7,] 0.10147146 0.005915522 0.03995173 -0.22876986
0.673039912 -0.39571449 0.571107765
```

**The next step would be calculating the PCs using the `prcomp()` function in R. Note that we will use the raw (uncentered) version of the data and have to specify the `center=TRUE` option to ensure the column means are trivial. We can save the model information into `pca1` where `pca1$rotation` provides the loadings for each PC.**

```
> pca1<-prcomp(as.matrix(ReserveData.sub), center = T)
> summary(pca1)
```

Importance of components:

	PC1	PC2	PC3	PC4
PC5				
PC6				
PC7				
Standard deviation	2541.5292	284.69447	139.79773	74.27570
	37.90522	27.31390	1.052	
Proportion of Variance	0.9835	0.01234	0.00298	0.00084
	0.00022	0.00011	0.000	
Cumulative Proportion	0.9835	0.99585	0.99883	0.99967
	0.99989	1.00000	1.000	

```
> pca1$rotation
```

	PC1	PC2	PC3	PC4
PC5	PC6	PC7		
SAVINGSL	0.43802915	0.840194054	0.15824598	0.22991060 -
	0.154725598	-0.01831078	0.005333293	
M2SL	0.86861886	-0.366656200	0.02497568	-0.32952256
	0.008871566	-0.04212471	-0.001550690	
M1NS	0.17100350	-0.373093272	-0.03754267	0.84575739 -
	0.331941670	0.06604044	-0.017250689	
BOGAMBSL	0.11643248	0.097111962	-0.60687967	0.17742808
	0.488043646	-0.06809777	-0.578277773	
TRARR	0.01957359	0.071528131	-0.64779682	-0.02767588 -
	0.186822349	-0.44733754	0.582306156	
BORROW	0.01238692	0.076528374	-0.42823496	-0.19664207 -
	0.373796867	0.79510424	-0.004673683	
CURRCIR	0.10147146	0.005915522	0.03995173	0.22876986
	0.673039912	0.39571449	0.571107765	

**We notice that the loadings are just the eigenvectors times -1. These loadings represent a vector in 6D space (we have 6 columns in the original data). The scale factor -1 just represents the opposite direction of the eigenvector. We can also load the factoextra package and compute the eigenvalues of each PC.**

```
> library(factoextra)
> eigen<-get_eigenvalue(pca1)
> eigen
```

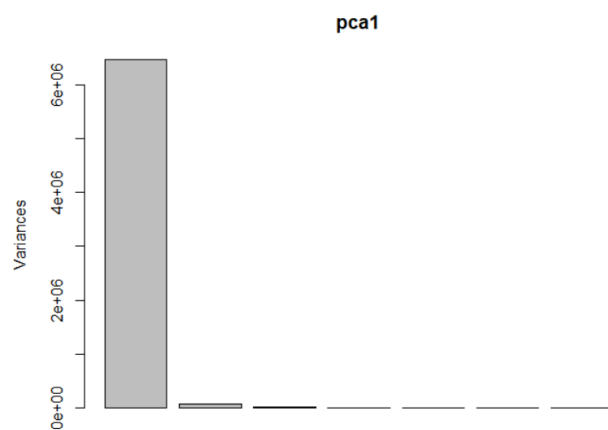
	eigenvalue	variance.percent	cumulative.variance.percent
Dim.1	6.459371e+06	9.835109e+01	98.35109
Dim.2	8.105094e+04	1.234090e+00	99.58518
Dim.3	1.954341e+04	2.975700e-01	99.88275



Dim.4	5.516879e+03	8.400060e-02
	99.96675	
Dim.5	1.436806e+03	2.187696e-02
	99.98862	
Dim.6	7.460492e+02	1.135943e-02
	99.99998	
Dim.7	1.106305e+00	1.684473e-05
	100.00000	


**To see a detailed information about the variances that each PC explain we utilize the `plot()` function to visualize the PC loadings.**

```
>plot(pca1)
```



```
>library(graphics)
```

```
>biplot(pca1, choices = 1:2, scale = 1, pc.biplot = F)
```



```
library("factoextra")
```

### Data for the supplementary qualitative variables

```
> qualit_vars <- as.factor(ReserveData.sub$CURRCIR)
> head(qualit_vars)
> fviz_pca_biplot(pca1, axes = c(1, 2), geom = c("point", "text"),
                  col.ind = "black", col.var = "steelblue", label = "all",
                  invisible = "none", repel = T, habillage =
qualit_vars,
                  palette = NULL, addEllipses = TRUE, title = "PCA -
Biplot")
```

```
> install.packages("tsne")
> library(tsne)
> install.packages("Rtsne")
> library(Rtsne)
> dim(ReserveData)
[1] 609 8
```

## # t-SNE

**Identify the label-nomenclature - digits 0, 1, 2, ..., 9 - and map to diff colors**

```
> ReserveData.borrow<-ReserveData$BORROW
> ReserveData$BORROW<-as.factor(ReserveData$BORROW)
> ReserveData.borrow.colors =
rainbow(length(unique(ReserveData$BORROW)))
> names(ReserveData.borrow.colors) =
unique(ReserveData$BORROW)
```

**May need to check and increase the RAM allocation**

```
> memory.limit()
[1] 4012
> memory.limit(50000)
[1] 50000
```

**Run the t-SNE, tracking the execution time (artificially reducing the sample-size to get reasonable calculation time)**

```
> execTime_tSNE <- system.time(tsne_digits <-
Rtsne(ReserveData[1:10000 , -1], dims = 2, perplexity=30,
verbose=TRUE, max_iter = 500))
```

Read the 609 x 50 data matrix successfully!

Using no\_dims = 2, perplexity = 30.000000, and theta = 0.500000

Computing input similarities...

Normalizing input...

Building tree...

- point 0 of 609

Done in 0.35 seconds (sparsity = 0.172109)!

Learning embedding...

Iteration 50: error is 54.090328 (50 iterations in 0.68 seconds)

Iteration 100: error is 47.739200 (50 iterations in 0.54 seconds)

Iteration 150: error is 46.353491 (50 iterations in 0.55 seconds)

Iteration 200: error is 45.691023 (50 iterations in 0.56 seconds)

Iteration 250: error is 45.212232 (50 iterations in 0.60 seconds)

Iteration 300: error is 0.268417 (50 iterations in 0.62 seconds)

Iteration 350: error is 0.188355 (50 iterations in 0.75 seconds)

Iteration 400: error is 0.166797 (50 iterations in 0.57 seconds)

Iteration 450: error is 0.157891 (50 iterations in 0.57 seconds)

Iteration 500: error is 0.151636 (50 iterations in 0.57 seconds)

Fitting performed in 6.02 seconds.

```
> execTime_tSNE
```

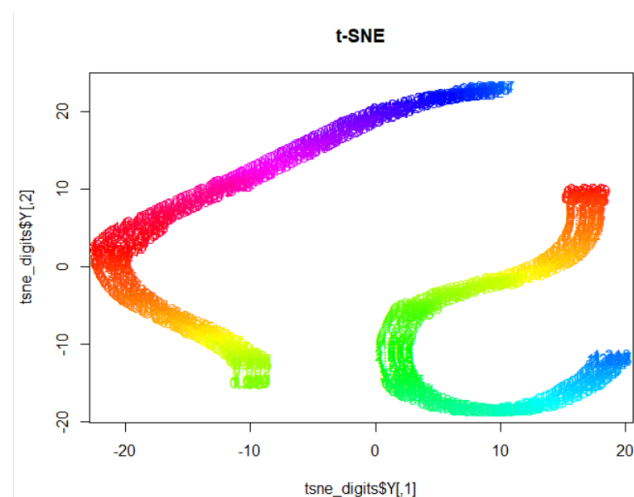
user	system	elapsed
------	--------	---------

8.23	0.08	11.17
------	------	-------

```
> plot(tsne_digits$Y, t='n', main="t-SNE") # don't plot the points  
to avoid clutter
```

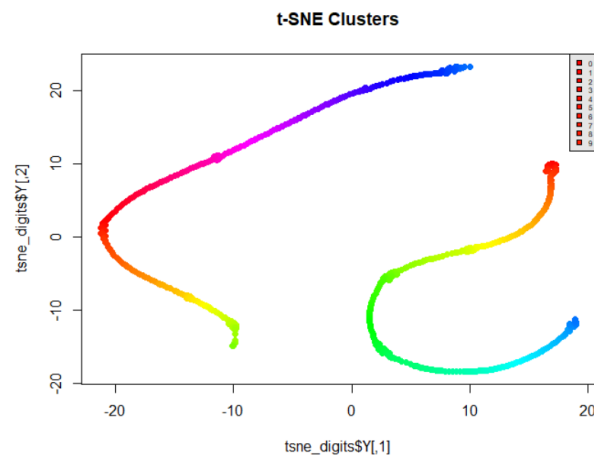
```
> text(tsne_digits$Y, labels=names(ReserveData.borrow.colors),  
col=ReserveData.borrow.colors)
```

```
> Y
```



```
>plot(tsne_digits$Y, main="t-SNE Clusters",
col=ReservedData.borrow.colors, pch = 19)

>legend("topright", c("0", "1", "2", "3", "4", "5", "6", "7", "8", "9"),
fill=ReservedData.borrow.colors, bg='gray90', cex=0.5)
```



### Problem 3.4 (Least Squares Estimation)

**Use the SOCR Knee Pain datasetLinks to an external site., extract the RB = Right-Back locations (x,y), and fit in a linear model for vertical location (y) in terms of the horizontal location (x). Display the linear model on top of the scatter plot of the paired data.**

#### # Load the SOCR Knee Pain Data

```
>wiki_url =
read_html("http://wiki.socr.umich.edu/index.php/SOCR_Data_Knee
PainData_041409")

>html_nodes(wiki_url, "#content")

>KneePain = html_table(html_nodes(wiki_url, "table")[[2]])

>KneePainData = as.data.frame(KneePain)

>KneePainData_sub = subset(KneePainData, View = RB)

>x=KneePainData_sub$x

>y=KneePainData_sub$Y>X <- cbind(1, x)

>beta_hat <- solve( t(X) %*% X ) %*% t(X) %*% y
```

```
###or
```

```
>beta_hat <- solve( crossprod(X) ) %*% crossprod( X, y )
```

**# Now we can see the results of this by computing the estimated  $\beta^0 + \beta^1 x$  for any value of  $x$ :**

```
>newx <- seq(min(x), max(x), len=100)
```

```
>X <- cbind(1, newx)
```

```
>fitted <- X%*%beta_hat
```

```
>plot(x, y)
```

```
>lines(newx, fitted, col=2)
```

